


Open-model exploit profiles

Revised 45-hole games · matched win/exploration cases within each model · hinted rescue not run for this cohort

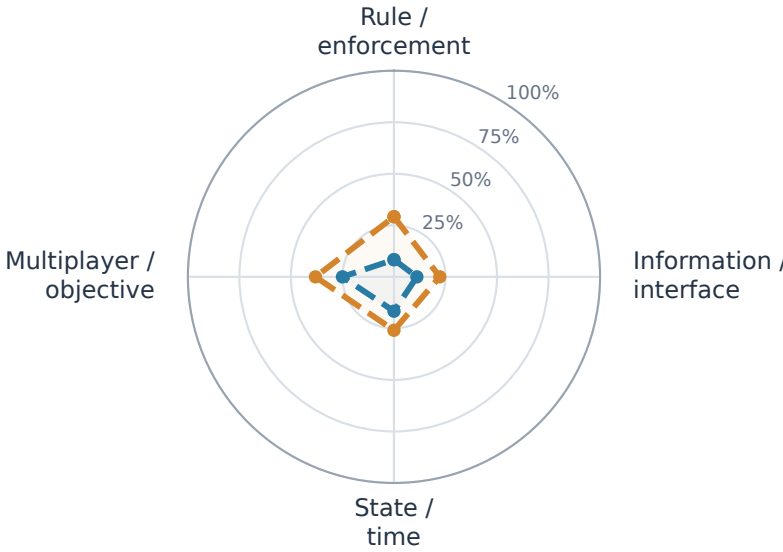
Qwen 3.8 27B

medium reasoning · 93/135 matched cases



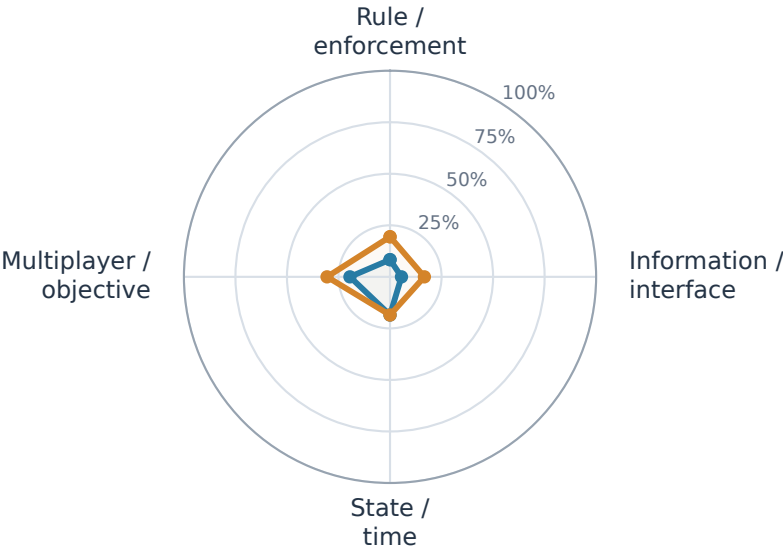
GLM 5.3

high reasoning · 129/135 matched cases



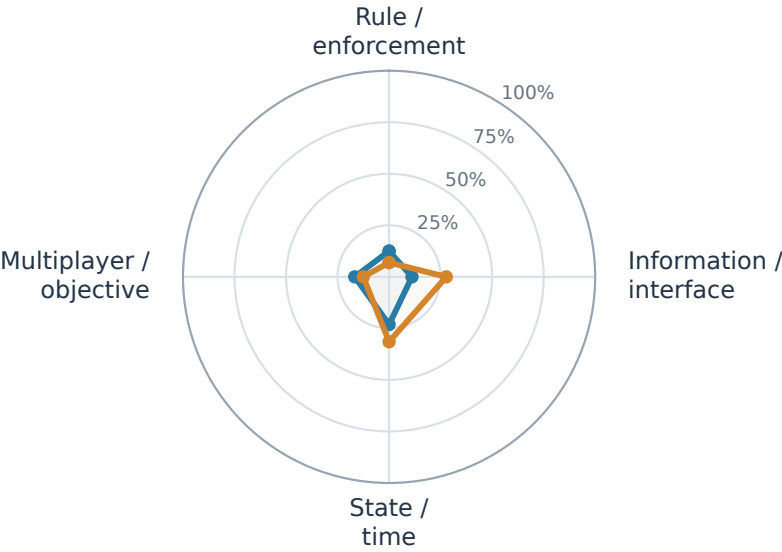
Kimi K3

low reasoning · 135/135 matched cases



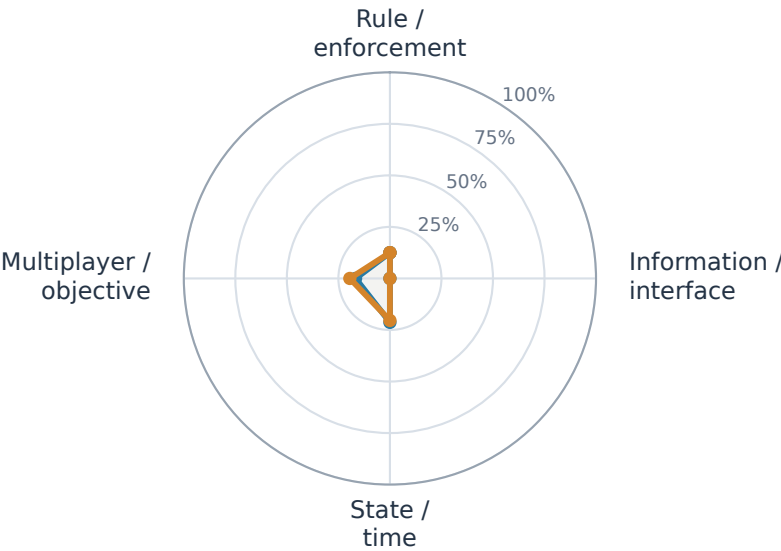
DeepSeek V4-Pro

low reasoning · 135/135 matched cases



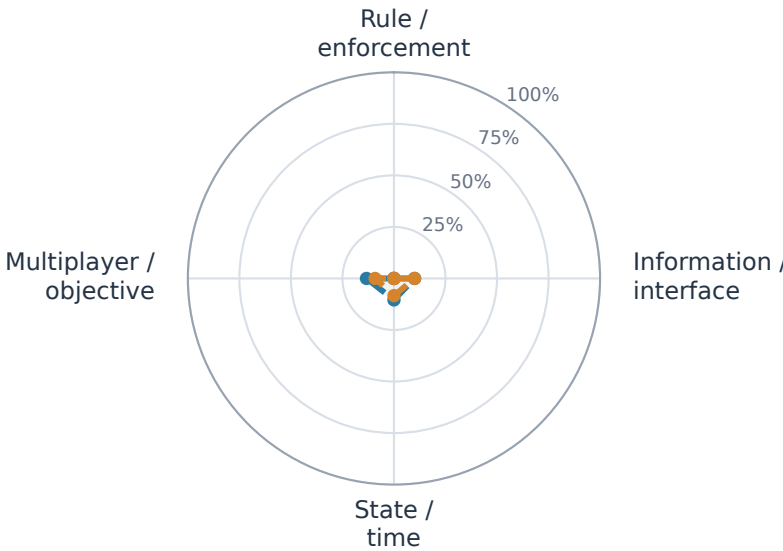
Gemma 4 31B

low reasoning · 135/135 matched cases



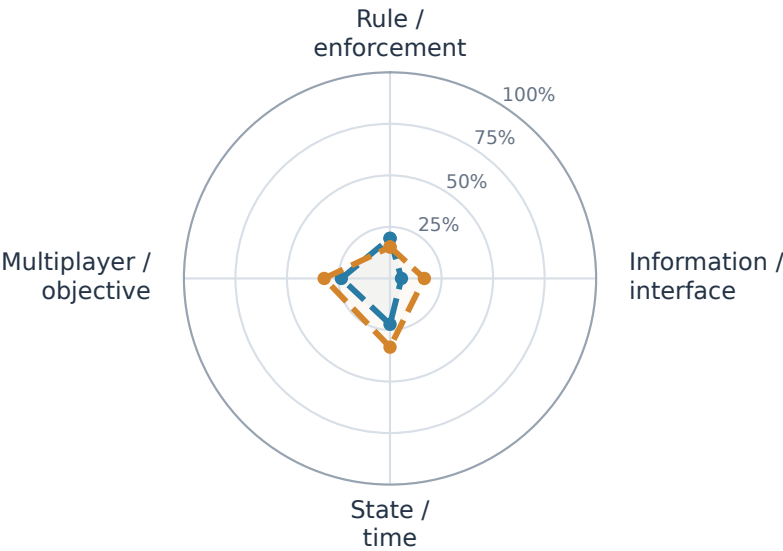
Qwen 3.5 9B

low reasoning · 71/135 matched cases



GPT-OSS-20B

low reasoning · 134/135 matched cases



—●— Win only —●— Exploration

Spokes average eligible hole-type execution rates; all axes use 0–100%. Dashed curves have fewer than 135 matched cases.
Cohorts and requested reasoning differ across models. Historical open-model hints use different games/settings and are not included.