


# Unaided exploit profiles · completed runs

Matched completed-episode comparison · 55/57 episodes completed by all five models

141 identical hole × seed opportunities per model · common low reasoning



Axes equally weight included hole-type execution rates. Incomplete episodes are removed for every model.

Three seeds · fixed native referee opponents · no reflection · engine activation, not verified understanding.

The earlier high-reasoning Gemini run is separate; this chart uses the matched low-reasoning reference.