



Model separation in v3-MA

Original four-model screen · ordinary versus privately nerfed recipients · no training or reflection

Ordinary recipients



Nerfed recipients



Haiku 4.5 GPT-5 mini Qwen 3.8 27B GLM 5.3

Macro-average of scenario marker rates within each family; 313 matched completed pairs.
Observed behavior, not inferred discovery. Commons excluded: no nerfed condition was sampled.