



Training depth versus few-shot prediction

Win / full-solution Brier (lower is better)

Higher parameter values · 160 episodes

Four unseen families · 112 episodes



- Learned linear
- 4-shot
- 4-shot + learned correction
- 8-shot
- 16-shot
- Pooled 4-shot (secondary)
- Pooled 16-shot (secondary)
- Estimated within-condition variance

Primary curves use fixed prospective tests; pooled controls were added later. Variance is a noisy floor component, not a proven bound.